



# Microservices with Micronaut

2021-01-28 - Moritz Kammerer

# About me

- 2013 M.Sc. Computer Science @ Munich University of Applied Sciences
- 2013 Software Engineer @ QAware
- ...
- 2021 Expert Software Engineer @ QAware

Working on backend stuff, mostly Java.

[LinkedIn](#) | [GitHub](#)



# Why Cloud?

# What's the problem?

```
Tomcat started on port(s): 8080 (http) with context path '/user'  
: Started UserManagementApplication in 13.307 seconds (JVM running for 13.856)  
: Git commit hash a3160f2d39df425a7f095fe4b07d412c6d2836e3, local changes: true  
: Initializing Spring FrameworkServlet 'dispatcherServlet'
```

```
Tomcat started on port(s): 3122 (http) with context path '/cvi'  
Started CviCoreApplication in 27.24 seconds (JVM running for 28.307)
```

# Why does it take so long?

- Start time of Spring (Boot)
  - Depends on the number of beans in the context
  - Reading of bytecode [3]
  - Creating proxies at runtime
  - Runtime AOP
  - Reflection
  - Annotation Synthesizing [1] [2]
- Memory usage
  - Reflective Metadata Cache
- Debugging
  - Very very long stacktraces
  - ... with code in dynamically generated proxy classes

[1] <https://docs.spring.io/spring-framework/docs/current/javadoc-api/org/springframework/core/annotation/AnnotationUtils.html>

[2] <https://docs.spring.io/spring-framework/docs/current/javadoc-api/org/springframework/core/annotation/AliasFor.html>

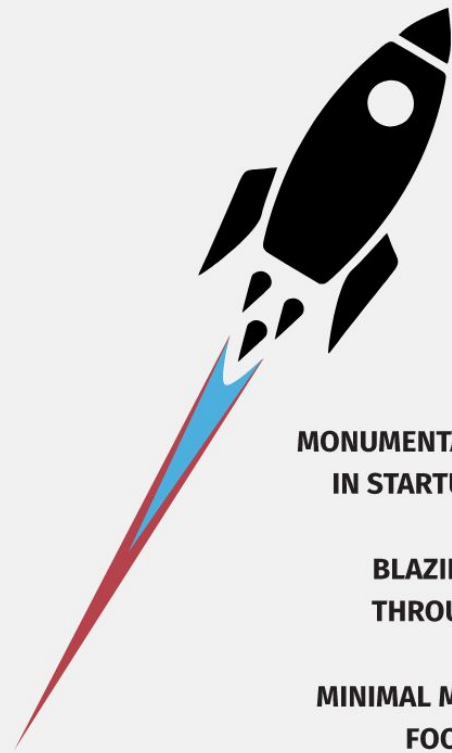
[3] <https://github.com/spring-projects/spring-framework/blob/master/spring-core/src/main/java/org/springframework/core/type/classreading/AbstractRecursiveAnnotationVisitor.java>

# Micronaut to the rescue!



M I C R O N A U T™

A modern, JVM-based, full-stack framework for building modular, easily testable microservice and serverless applications.



MONUMENTAL LEAP  
IN STARTUP TIME

BLAZING FAST  
THROUGHPUT

MINIMAL MEMORY  
FOOTPRINT















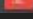
# What's the (promised) solution?

- Create stuff at compile time
  - Proxies
  - Dependency Injection
  - AOP
- Why?
  - Compiled once, run often
  - Tradeoff: Compile time vs. Runtime
- Goal:
  - No reflection - or as little as possible
  - No proxies
  - Minimize startup time
  - Minimize memory usage
  - Readable stacktraces

# And how does that work?

- Annotation Processor
  - Reads bytecode & annotations of the classes
  - Creates new classes
  - Doesn't touch your code, just extends it
- Gradle
  - `annotationProcessor "io.micronaut:micronaut-inject-java"`
  - `annotationProcessor "io.micronaut:micronaut-validation"`
- Also runs with Maven
  - `< Insert 500 KB of XML here >`



-  \$CacheWarmerDefinition.class
-  \$CacheWarmerDefinition\$\$exec1.class
-  \$CacheWarmerDefinition\$\$exec1\$\$AnnotationMetadata.class
-  \$CacheWarmerDefinitionClass.class
-  \$CacheWarmerDefinitionClass\$\$AnnotationMetadata.class
-  \$ProjectsControllerDefinition.class
-  \$ProjectsControllerDefinition\$\$exec1.class
-  \$ProjectsControllerDefinition\$\$exec1\$\$AnnotationMetadata.class
-  \$ProjectsControllerDefinitionClass.class
-  \$ProjectsControllerDefinitionClass\$\$AnnotationMetadata.class
-  Application.class
-  CacheWarmer.class
-  MicronautSubstitutions.class
-  ProjectDto.class
-  ProjectsController.class

# And this is really working?

- Spring Boot 2.1.2 (starter-web), Java 11, @Component on empty class
  - 0 Beans: 1.236 seconds
  - 1000 Beans: 1.808 seconds
  - 2000 Beans: 2.356 seconds
  - 4000 Beans: 3.259 seconds
  - 8000 Beans: 6.498 seconds
  - 16000 Beans: 9.231 seconds
- Micronaut 1.0.4, Java 11, @Singleton on empty class
  - 0 Beans: 1085ms
  - 1000 Beans: 1870ms
  - 2000 Beans: 2757ms
  - 4000 Beans: 4529ms
  - 8000 Beans: 8038ms
  - 16000 Beans: 15861ms

Hmmm...



# Debunked: It's not the classpath scanning

- Not so expensive:
  - Classpath scanning (Spring)
  - Injection with reflection (Spring)
- Expensive
  - Class Loading (Micronaut + Spring)
  - Runtime AOP Proxies (Spring)
  - Annotation Synthesizing (Spring)
  - Read bytecode on big classes (Spring)
  - Auto-Configuration Magic (Spring)
- Most of the expensive stuff is done at Micronaut compile time

# Looks great, so not disadvantages?



**Jochen Mader** 

@codepitbull

Folgen



A good developer is like a werewolf: Afraid of silver bullets.

02:48 - 8. Okt. 2016

437 Retweets 524 „Gefällt mir“-Angaben



10



437



524

# The price you pay: Compile time (16k Beans)

Spring:

[INFO] Total time: 7.958 s

Micronaut:

BUILD SUCCESSFUL in 5m 3s



# Benchmarking is hard!

- This was **NOT** a real benchmark
- Most time spent in the applications shown on the first slides: Spring autoconfiguration magic
- It really depends on the needed feature of your application
  - Actuator
  - JPA / Hibernate
  - Tracing
  - Metriken
  - @FeignClient
  - @Transactional
  - @Cacheable
  - ...
- More in the live demo!

# Okay... what else does Micronaut bring to the table?

- Dependency Injection
  - @Inject, @Named, @Singleton, @Prototype, ...
- REST Controller
  - @Controller, @GET, @POST, ...
- Events
  - @PostConstruct
  - @EventListener
- AOP
  - Around Advice, Introduction Advice, ...
- Validation
  - @Validated, @NotNull, ...



...

- Caching
  - @Cachable, @CacheInvalidate, ...
- Retry
  - @Retryable
- Circuit Breaking
  - @CircuitBreaker
- Fallback / Recovery
  - @Recoverable / @Fallback
- Scheduling
  - @Scheduled

...

- Reactive & Blocking HTTP with Netty
  - Compile Time Route Validation
  - @Error Exception Handling ala Spring
- Websocket Server & Clients
- Server side views
  - Thymeleaf, Handlebars, Velocity
- OpenAPI / Swagger
  - Open API spec is created on compile time!
- HTTP Clients
  - @Client
  - Client-side load balancing
  - Service discovery (Consul, Eureka, K8S, Route 53, DNS based)
  - OpenTracing (Zipkin, Jaeger)

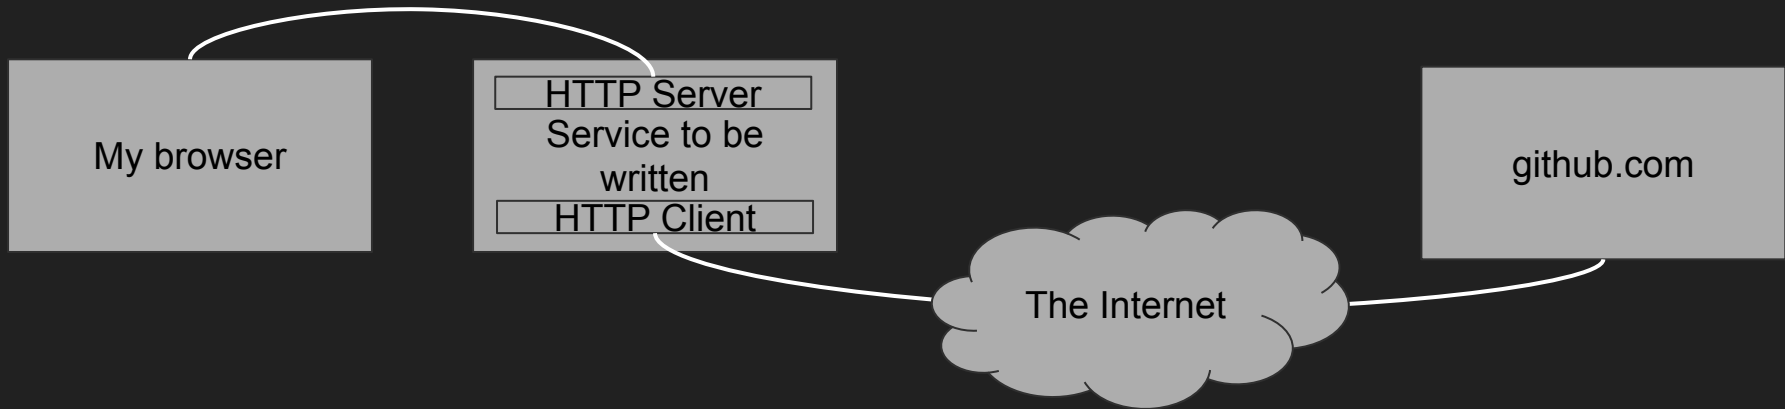
...

- Application Configuration
  - application.yml / .json / .groovy
  - Different Environments
  - Property Sources (ENV, System Properties, ...)
  - @ConfigurationProperties, @Value
  - Distributed config with Consul, AWS Parameter Store
- Database support
  - JPA, Mongo, Neo4J, Postgres-Reactive, Redis, Cassandra
- Monitoring
  - @Endpoint (/beans, /info, /health, ...)
  - Metriken with Micrometer
  - Change log level at runtime

...

- Security
  - Like Spring Security with users and roles (RBAC)
  - Basic Auth / Session based
  - JWT (JWS and JWE) incl. automatic token propagation
  - Auth with LDAP
- Function-as-a-Service (AWS Lambda, OpenFaaS)
- Message driven services (RabbitMQ / Kafka)
- CLI Applikationen
- Good documentation!
  - e.g. with tutorials for Let's Encrypt, starting Consul, etc.

Enough bla bla, show code!



# Benchmarks

# Spring vs. Micronaut - GitHub Scraper

## **Micronaut 2.2.1**

BUILD SUCCESSFUL in 3s

0,89s user 0,07s system 22% cpu 4,199 total

Startup completed in 625ms



# Spring vs. Micronaut - GitHub Scraper

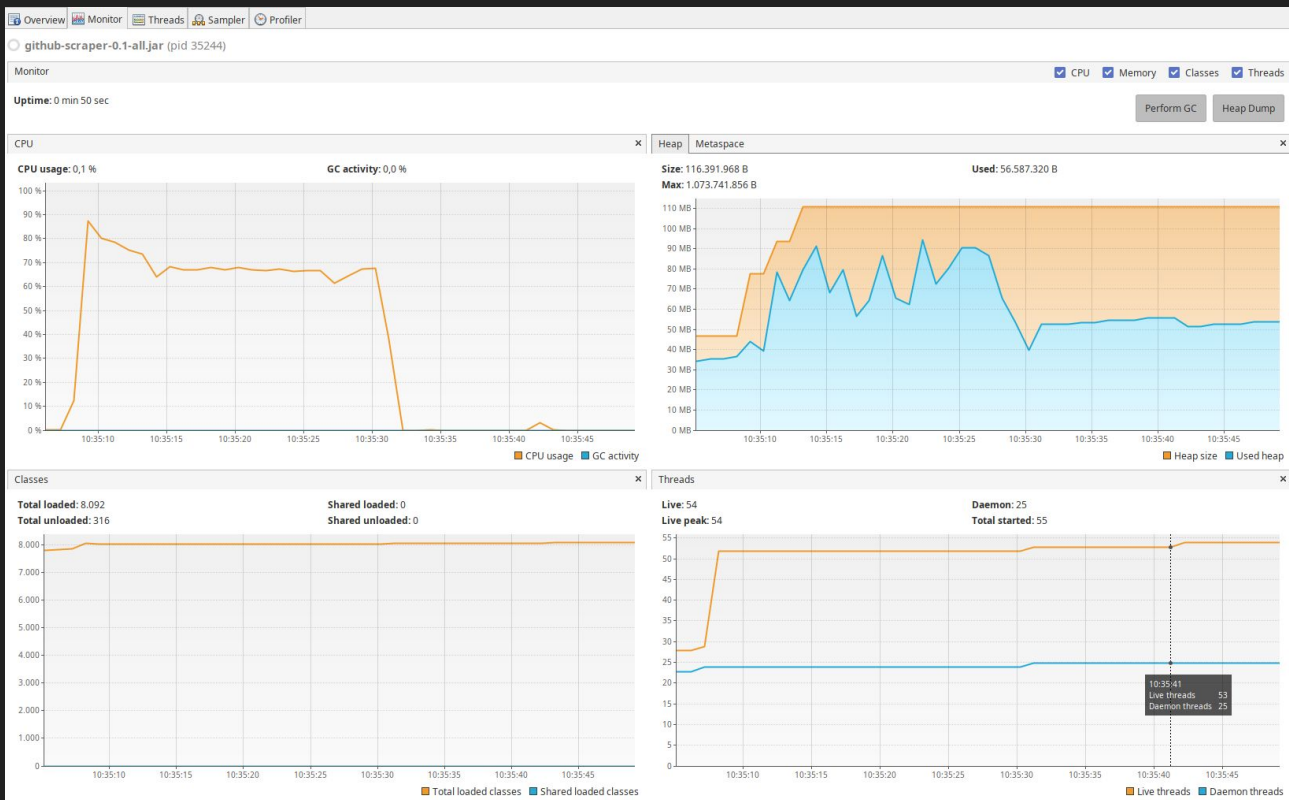
## Spring 2.4.1

BUILD SUCCESSFUL in 1s

0,86s user 0,09s system 67% cpu 1,414 total

Started DemoApplication in 1.21 seconds (JVM running for 1.556)

# Micronaut (JAR: 14.8 MB)



```
java -Xms16M -Xmx512M -jar github-scraper-0.1-all.jar / OpenJDK Runtime Environment AdoptOpenJDK (build 11.0.9+11)
```

# Micronaut

```
wrk -t 8 -c 80 -d 30s --latency http://localhost:8080/github/stars
```

```
Running 30s test @ http://localhost:8080/github/stars
```

```
8 threads and 80 connections
```

Thread Stats	Avg	Stdev	Max	+/-	Stdev
--------------	-----	-------	-----	-----	-------

Latency	2.21ms	5.33ms	139.33ms	91.82%	
---------	--------	--------	----------	--------	--

Req/Sec	15.22k	5.34k	27.39k	72.74%	
---------	--------	-------	--------	--------	--

```
Latency Distribution
```

```
50% 286.00us
```

```
75% 1.98ms
```

```
90% 6.44ms
```

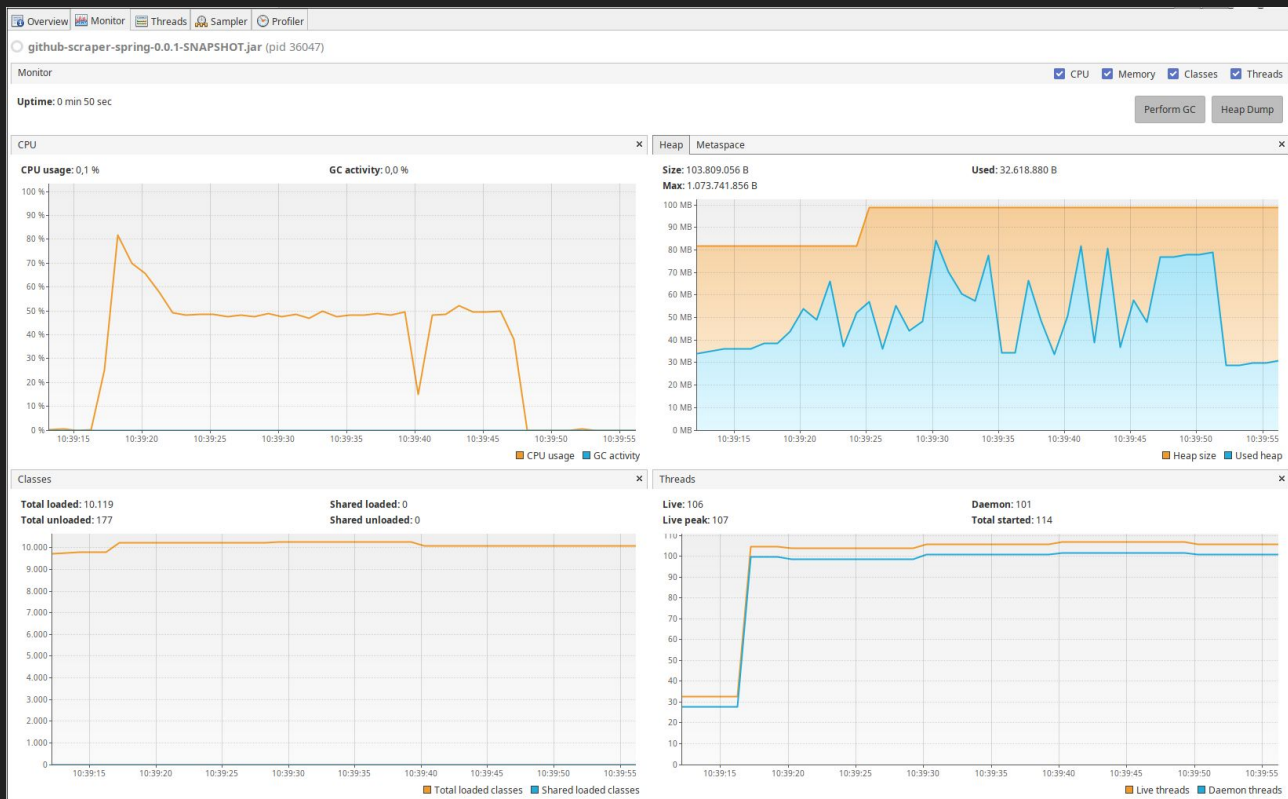
```
99% 22.34ms
```

```
2765210 requests in 30.07s, 1.38GB read
```

```
Requests/sec: 91956.68
```

```
Transfer/sec: 47.01MB
```

# Spring Boot (JAR: 30.6 MB)



```
java -Xms16M -Xmx512M -jar github-scraper-spring-0.0.1-SNAPSHOT.jar / OpenJDK Runtime Environment AdoptOpenJDK (build 11.0.9+11)
```

# Spring Boot

```
wrk -t 8 -c 80 -d 30s --latency http://localhost:8080/github/stars
```

```
Running 30s test @ http://localhost:8080/github/stars
```

```
8 threads and 80 connections
```

Thread Stats	Avg	Stdev	Max	+/-	Stdev
--------------	-----	-------	-----	-----	-------

Latency	10.23ms	65.15ms	778.07ms	97.82%	
---------	---------	---------	----------	--------	--

Req/Sec	9.29k	2.69k	14.24k	78.93%	
---------	-------	-------	--------	--------	--

```
Latency Distribution
```

```
50% 0.87ms
```

```
75% 1.55ms
```

```
90% 3.36ms
```

```
99% 414.50ms
```

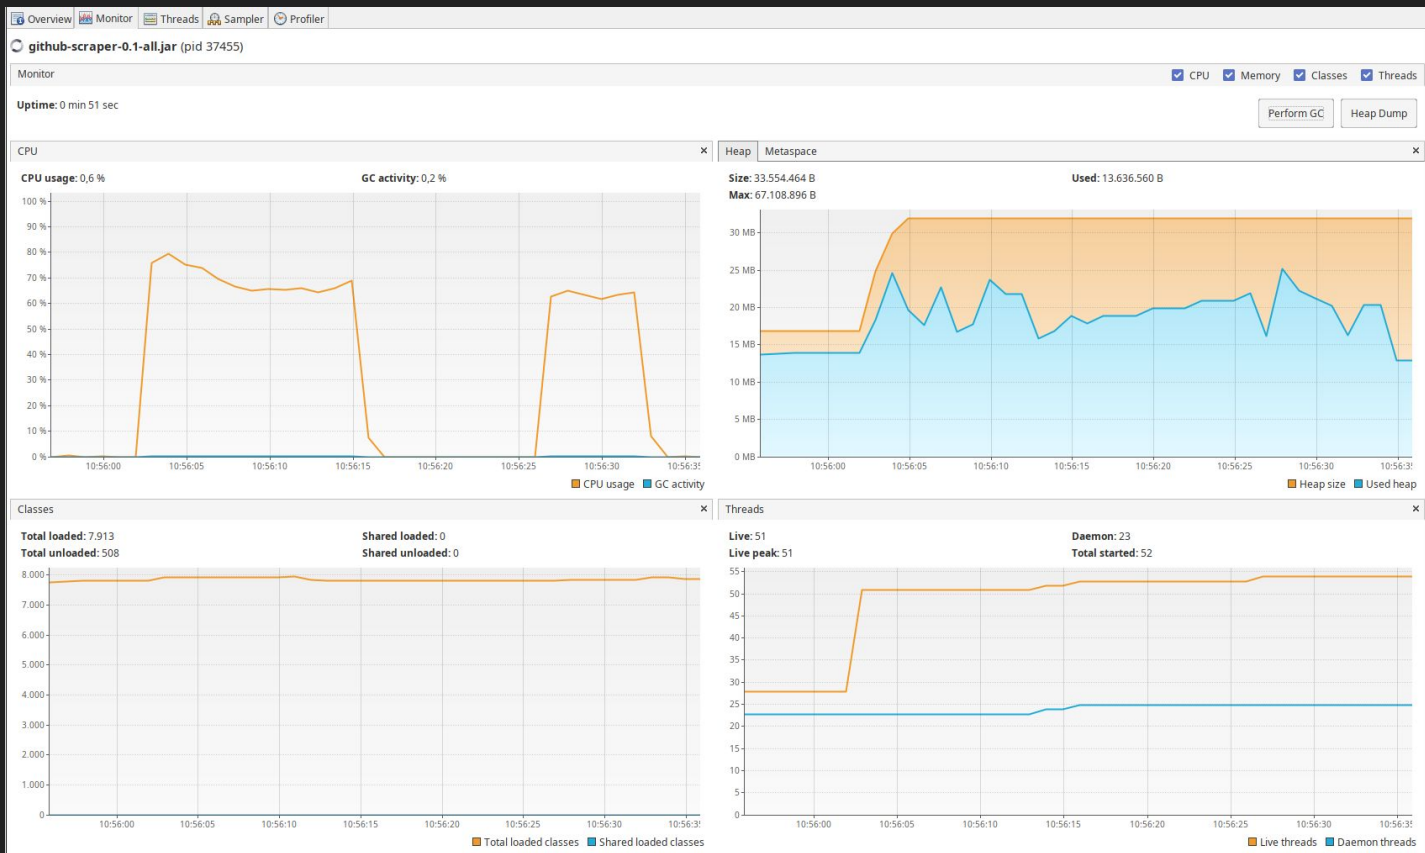
```
2169903 requests in 30.10s, 1.07GB read
```

```
Requests/sec: 72091.42
```

```
Transfer/sec: 36.38MB
```

How much RAM do I have to use?

# Minimal Heap Size - Micronaut (-Xmx32M)



# Minimal Heap Size - Micronaut (-Xmx32M)

```
wrk -t 8 -c 80 -d 30s --latency http://localhost:8080/github/stars
```

```
Running 30s test @ http://localhost:8080/github/stars
```

```
8 threads and 80 connections
```

Thread Stats	Avg	Stdev	Max	+/-	Stdev
--------------	-----	-------	-----	-----	-------

Latency	2.68ms	6.40ms	135.54ms	91.54%	
---------	--------	--------	----------	--------	--

Req/Sec	11.34k	4.58k	23.37k	68.22%	
---------	--------	-------	--------	--------	--

```
Latency Distribution
```

```
50% 270.00us
```

```
75% 2.58ms
```

```
90% 8.02ms
```

```
99% 27.00ms
```

```
1717518 requests in 30.06s, 631.95MB read
```

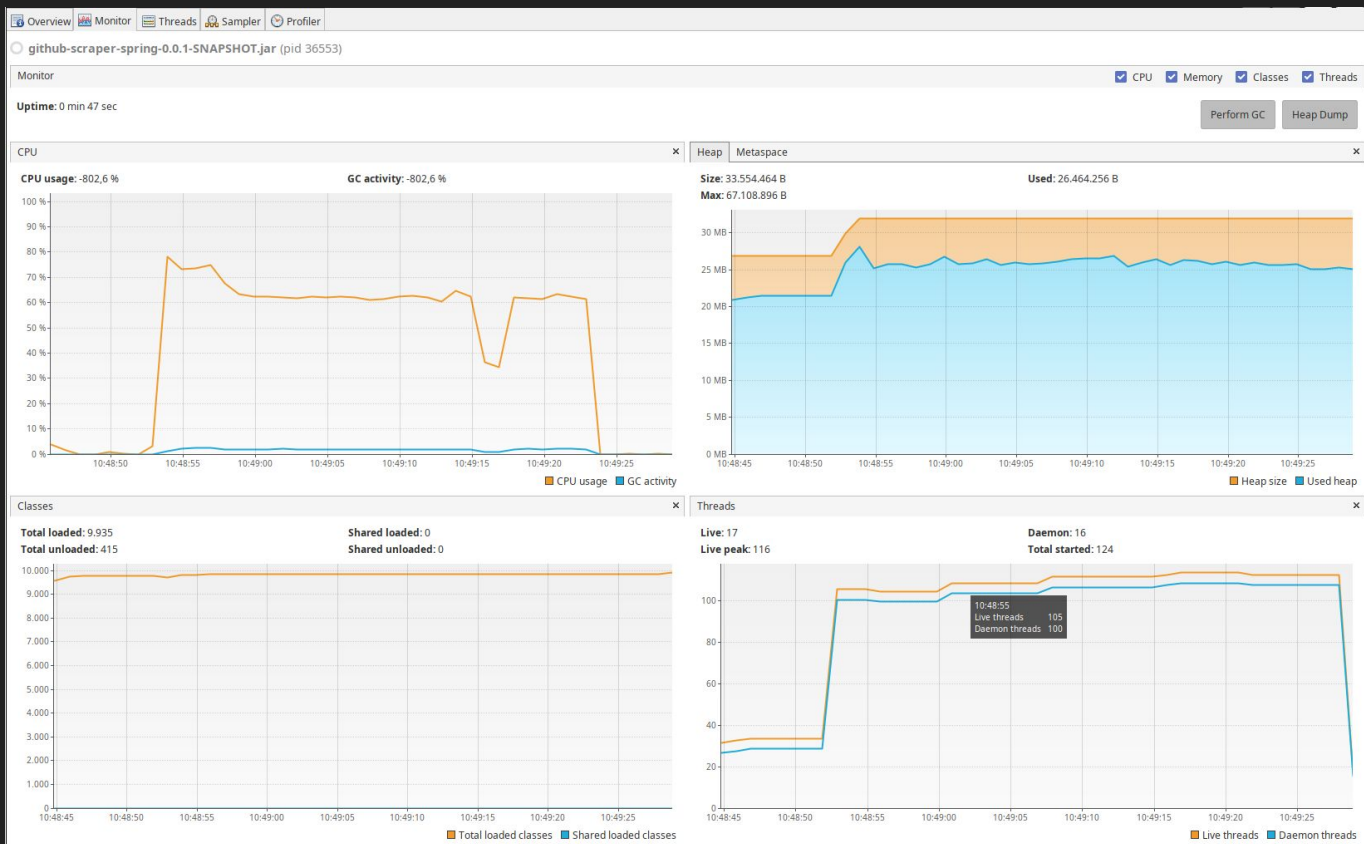
```
Socket errors: connect 0, read 0, write 0, timeout 80
```

```
Requests/sec: 57134.45
```

```
Transfer/sec: 21.02MB
```



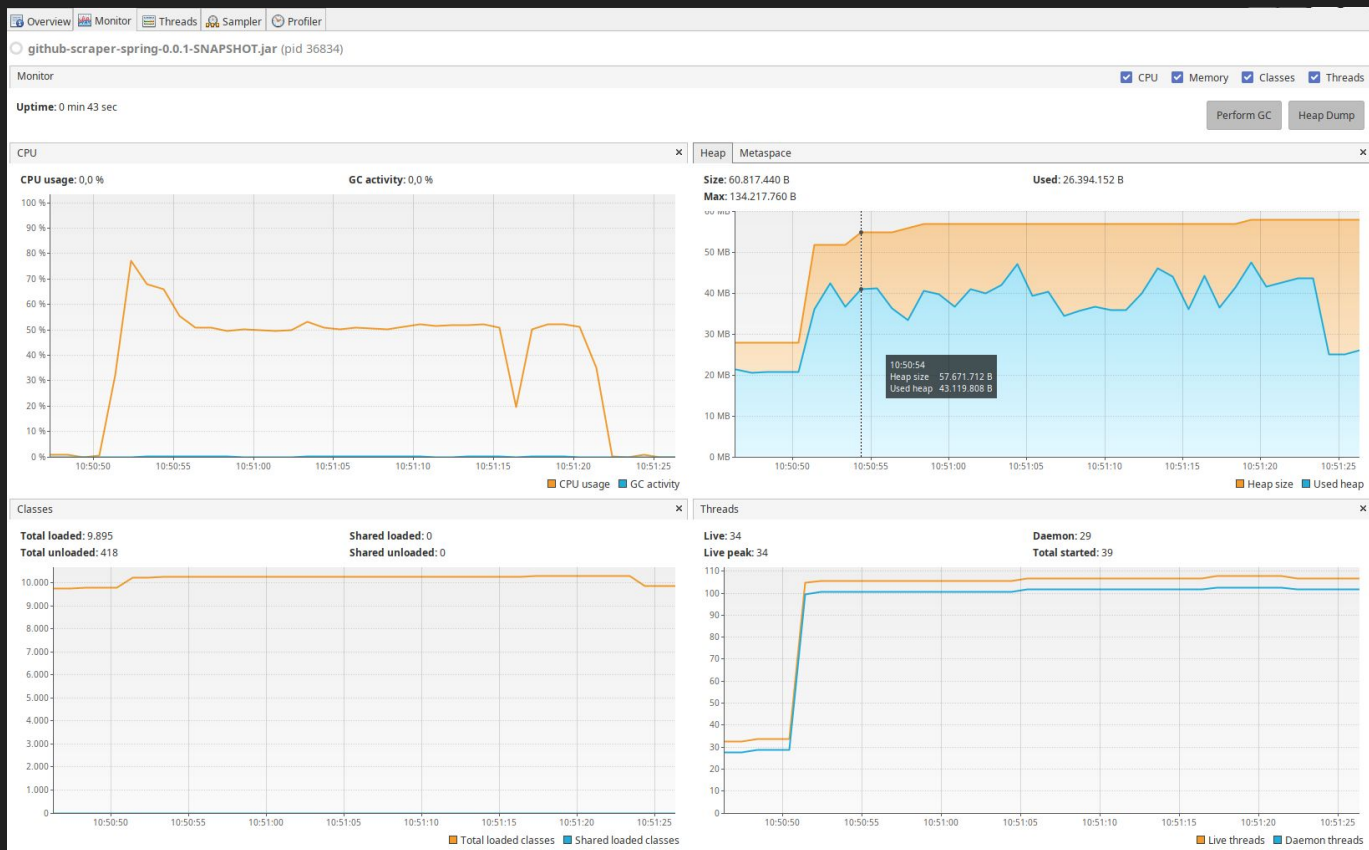
# Minimal Heap Size - Spring (-Xmx32M)



# Minimal Heap Size - Spring (-Xmx32M)

```
wrk -t 8 -c 80 -d 30s --latency http://localhost:8080/github/stars
Running 30s test @ http://localhost:8080/github/stars
 8 threads and 80 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency 16.38ms   87.16ms 955.97ms   97.33%
  Req/Sec 3.92k    1.19k   6.21k    75.93%
Latency Distribution
 50% 2.11ms
 75% 4.21ms
 90% 8.59ms
 99% 578.98ms
905844 requests in 30.05s, 457.16MB read
Requests/sec: 30148.82
Transfer/sec: 15.22MB
```

# Minimal Heap Size - Spring (-Xmx64M)



# Minimal Heap Size - Spring (-Xmx64M)

```
wrk -t 8 -c 80 -d 30s --latency http://localhost:8080/github/stars
```

```
Running 30s test @ http://localhost:8080/github/stars
```

```
8 threads and 80 connections
```

Thread Stats	Avg	Stdev	Max	+/-	Stdev
--------------	-----	-------	-----	-----	-------

Latency	10.50ms	66.17ms	786.18ms	97.73%	
---------	---------	---------	----------	--------	--

Req/Sec	8.60k	2.27k	13.25k	79.23%	
---------	-------	-------	--------	--------	--

```
Latency Distribution
```

```
50% 0.95ms
```

```
75% 1.68ms
```

```
90% 3.30ms
```

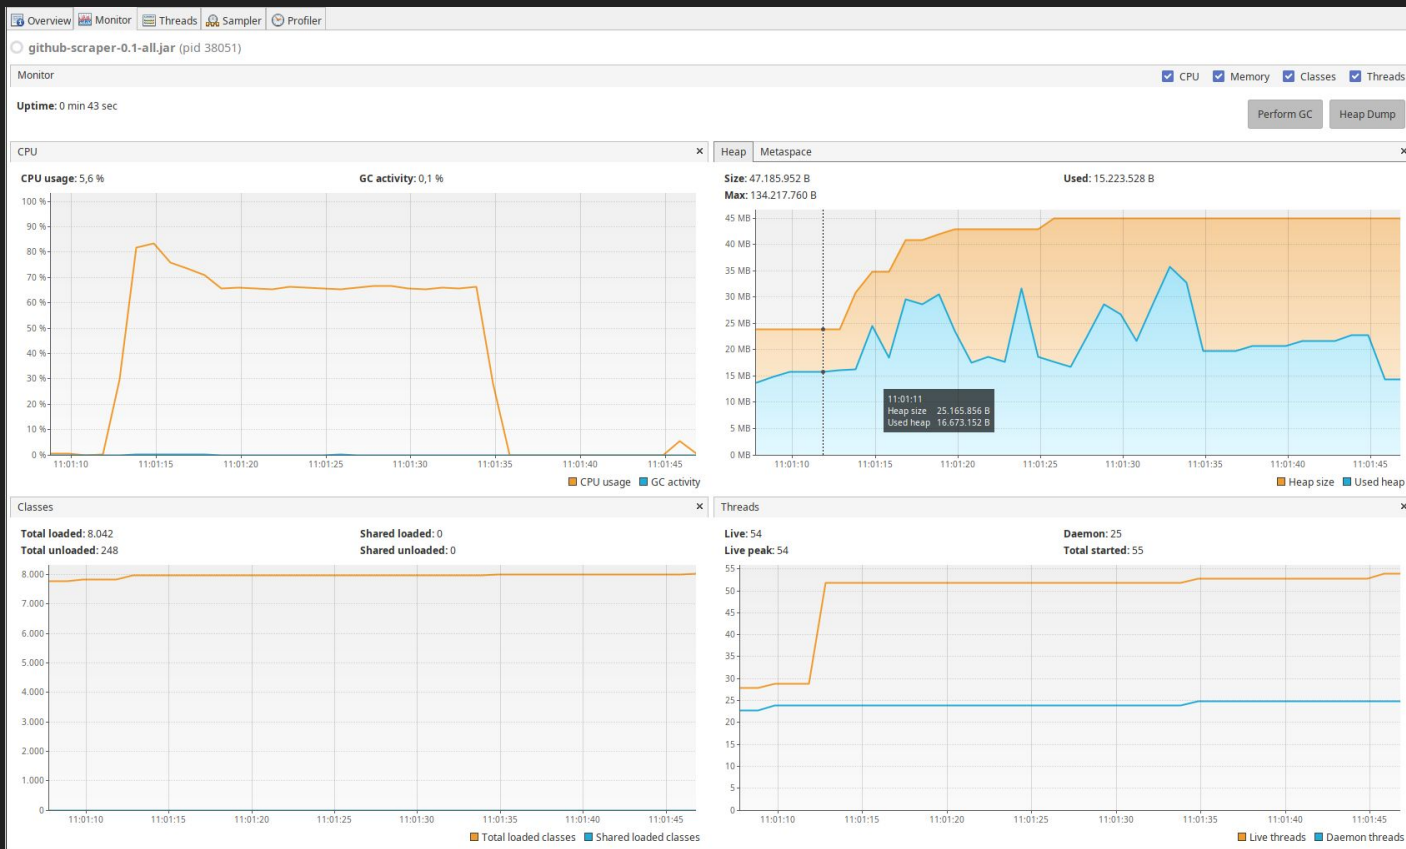
```
99% 423.31ms
```

```
2006502 requests in 30.09s, 0.99GB read
```

```
Requests/sec: 66689.38
```

```
Transfer/sec: 33.66MB
```

# Minimal Heap Size - Micronaut (-Xmx64M)



# Minimal Heap Size - Micronaut (-Xmx64M)

```
wrk -t 8 -c 80 -d 30s --latency http://localhost:8080/github/stars
```

```
Running 30s test @ http://localhost:8080/github/stars
```

```
8 threads and 80 connections
```

Thread Stats	Avg	Stdev	Max	+/-	Stdev
--------------	-----	-------	-----	-----	-------

Latency	2.28ms	5.05ms	95.77ms	90.92%	
---------	--------	--------	---------	--------	--

Req/Sec	14.11k	5.15k	25.09k	73.51%	
---------	--------	-------	--------	--------	--

```
Latency Distribution
```

```
50% 283.00us
```

```
75% 2.12ms
```

```
90% 6.78ms
```

```
99% 23.54ms
```

```
2451595 requests in 30.06s, 1.22GB read
```

```
Requests/sec: 81556.51
```

```
Transfer/sec: 41.69MB
```

# Graal VM



**“High-performance polyglot VM”**  
(and AOT compiler)

Enough bla bla, show code!



# Graal Substrate VM

```
$ ./gradlew nativeImage (or ./mvnw package -Dpackaging=native-image)
```

```
... 2 minutes later ...
```

```
$ file application
```

```
application: ELF 64-bit LSB executable, x86-64, version 1 (SYSV), dynamically linked,  
interpreter /lib64/ld-linux-x86-64.so.2 for GNU/Linux 3.2.0, with debug_info, not stripped
```

```
$ ll application
```

```
-rwxr-xr-x. 1 moe moe 62M 18. Dez 11:29 application
```

```
$ ldd application
```


```
linux-vdso.so.1 (0x00007ffd7cbbe000)  
libpthread.so.0 => /lib64/libpthread.so.0 (0x00007fc717414000)  
libdl.so.2 => /lib64/libdl.so.2 (0x00007fc71740d000)  
libz.so.1 => /lib64/libz.so.1 (0x00007fc7173f3000)  
librt.so.1 => /lib64/librt.so.1 (0x00007fc7173e8000)  
libgcc_s.so.1 => /lib64/libgcc_s.so.1 (0x00007fc7173cd000)  
libc.so.6 => /lib64/libc.so.6 (0x00007fc717202000)  
/lib64/ld-linux-x86-64.so.2 (0x00007fc717456000)
```

# Graal (no -Xmx set)

Startup completed in 67ms.

Idle:

PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME+	START	Command
57535	moe	20	0	1390M	94524	46008	S	0.0	0.3	0:00.08	13:33	./application



KB = 92 MB

Under load:

PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME+	START	Command
57535	moe	20	0	4135M	764M	49276	S	732.	2.4	1:42.99	13:33	./application

# Graal - Benchmark

```
wrk -t 8 -c 80 -d 30s --latency http://localhost:8080/github/stars
```

```
Running 30s test @ http://localhost:8080/github/stars
```

```
8 threads and 80 connections
```

Thread Stats	Avg	Stdev	Max	+/-	Stdev
--------------	-----	-------	-----	-----	-------

Latency	1.60ms	2.28ms	41.76ms	86.51%	
---------	--------	--------	---------	--------	--

Req/Sec	8.88k	1.21k	12.09k	79.71%	
---------	-------	-------	--------	--------	--

```
Latency Distribution
```

```
50% 416.00us
```

```
75% 2.34ms
```

```
90% 4.61ms
```

```
99% 10.16ms
```

```
1355757 requests in 30.04s, 305.29MB read
```

```
Socket errors: connect 0, read 0, write 0, timeout 80
```

```
Requests/sec: 45130.28
```

```
Transfer/sec: 10.16MB
```

# Conclusion

# Experience from the trenches

## Pro:

- Services start faster than Spring Boot
- Smaller JARs
- Tests are faster, as the embedded server starts fast
  - Spring mitigates this with application context caching!

## Contra:

- No separation of management and API port
- Spring Boot is more common (more documentation, blog posts, Stack Overflow answers)
- Extension story from Spring is better (e.g. OAuth2)

# Conclusion

- Good documentation
  - Recommendation for writing small services
    - I have no experience with bigger services in Micronaut. If you have, please contact me :)
  - Development is fun
  - If the features are sufficient: Try it!
- 
- GraalVM Substrate VM, if you get it to run, is really cool
    - This is NOT a JVM, so profilers / metrics / thread dumps etc. won't work
    - Build times like it's C++ (> 30s)
    - Not all libraries are compatible and workarounds must be employed
      - I ran into trouble with Caffeine Cache ([GitHub Issue](#))

# Literature

- GitHub Repo: <https://github.com/qaware/microservices-with-micronaut>
- Micronaut documentation: <https://docs.micronaut.io/latest/guide/index.html>
- GraalVM: <https://www.graalvm.org/>

# Questions?

Contact: [moritz.kammerer@qaware.de](mailto:moritz.kammerer@qaware.de)



# We are hiring.

In Munich. And in Mainz.

<https://www.qaware.de/karriere/#jobs>